

Noah French (njf5cu)

Lab 9

11/16/17

Filename: inlab9.pdf

File Description: Lab 9 In-lab Report

Optimized Code

To compare assembly code generated without and without the -O2 flag, I first created a simple C++ function that takes in an integer, and uses it to loop through zero and that integer. Every iteration of the loop, the function adds one to an integer that originally starts at 0. It returns the result. The function is, in actual practice, completely useless, because it effectively returns the same integer value that it takes in as a parameter. It should be useful, however, in analyzing how loops and function calls are represented on the assembly level. The function looks like this:

```
int loopFunc(int x) {
    int y = 0;
    for (int i=0; i < x; i++)
        y++;
    return y;
}
```

In the main method of the same .cpp file, I initialize an integer z to equal loopFunc(4).

```
int main() {
    int z = loopFunc(4);
    return 0;
}
```

First, I compiled the .cpp file using the following flags: -m64 -mllvm -x86-asm-syntax=intel -S -fomit-frame-pointer. This is how the loopFunc() function was represented in x86:

```
    mov     dword ptr [rsp - 4], edi
    mov     dword ptr [rsp - 8], 0
    mov     dword ptr [rsp - 12], 0
.LBB1_1:
    mov     eax, dword ptr [rsp - 12]
    cmp     eax, dword ptr [rsp - 4]
    jge     .LBB1_4
    mov     eax, dword ptr [rsp - 8]
    add     eax, 1
    mov     dword ptr [rsp - 8], eax
    mov     eax, dword ptr [rsp - 12]
    add     eax, 1
    mov     dword ptr [rsp - 12], eax
    jmp     .LBB1_1
.LBB1_4:
    mov     eax, dword ptr [rsp - 8]
    retn
```

The loop is clearly evident in the x86 code. Every iteration, the routine jumps back up to the .LBB1_1 marker. One is added to the eax register each iteration. And at the end, the value that had been growing every iteration is moved to the return register, where the value is returned. Here's how the same function was represented in x86 when the -O2 flag was also included:

```
xor     eax, eax
test    edi, edi
cmovns  eax, edi
ret
```

The optimizer flag cleverly recognized what my loop function was doing. Because the loop function ended up returning the input value every time, the assembly code skipped the loop entirely. Instead, it simply took the value from the parameter register, edi, and put it into the return register, edx.

The difference between the x86 code for the caller function (in this case, the main method), was similarly striking. Without the -O2 flag the main method looked like this:

```
push    rax
.Ltmp1:
.cfi_def_cfa_offset 16
mov     edi, 4
mov     dword ptr [rsp + 4], 0
call    _Z8loopFunci
xor     edi, edi
mov     dword ptr [rsp], eax
mov     eax, edi
pop     rcx
ret
```

The value 4 is moved into the parameter register, so it can be passed into the loopFunc(). The resulting value from the function call is stored in memory (int z). The edi register is promptly zeroed, and then that zero is moved to the return register, because the main method always returns 0. The x86 code is distinctly less interesting to analyze when the -O2 flag is used:

```
xor     eax, eax
ret
```

It's clearly evident that the -O2 flag did not deem my task of making `int z = loopFunc(4)` worthy, and (as per the in-lab instruction's warning) it was completely deleted. Presumably, initialize a value is not on its own necessary. These examples exhibit a few key ways that the -O2 flag optimizes code: for one, it looks at the input and output of a function. For one, if the body of the function can achieve the same input → output mapping in fewer steps, it will be rewritten in fewer steps. Second, any unused variables will be completely deleted.

For more compelling examples, I looked at the getNextPrime() method from lab 6 in x86. Without the -O2 flag, the x86 code looked like this:

```
push    rax
```

```
.Ltmp39:
.cfi_def_cfa_offset 16
```

```

        mov     dword ptr [rsp + 4], edi
.LBB11_1:      # =>This Inner Loop Header: Depth=1
        mov     eax, dword ptr [rsp + 4]
        add     eax, 1
        mov     dword ptr [rsp + 4], eax
        mov     edi, eax
        call    _Z10checkprimej
        xor     al, -1
        test    al, 1
        jne     .LBB11_2
        jmp     .LBB11_3
.LBB11_2:      # in Loop: Header=BB11_1 Depth=1
        jmp     .LBB11_1
.LBB11_3:
        mov     eax, dword ptr [rsp + 4]
        pop     rcx
        ret

```

With the -O2 flag, the code looked like this:

```

        mov     ecx, edi
        .align  16, 0x90
.LBB5_1:
        inc     ecx
        cmp     ecx, 2
        jb     .LBB5_1
        mov     eax, 2
        je     .LBB5_6
        test    cl, 1
        je     .LBB5_1
        mov     esi, 5
        cmp     ecx, 9
        jb     .LBB5_5
        .align  16, 0x90
.LBB5_8:
        lea     edi, [rsi - 2]
        xor     edx, edx
        mov     eax, ecx
        div     edi
        test    edx, edx
        je     .LBB5_1
        mov     eax, esi
        imul    eax, eax
        add     esi, 2
        cmp     eax, ecx
        jbe     .LBB5_8
.LBB5_5:
        mov     eax, ecx
.LBB5_6:      # %.loopexit

```

ret

Interestingly, this snippet of code actually got longer with the optimization! However, because massive number of times this function loops to generate prime numbers, even though the optimized code is more lines, it is still more efficient. At runtime, the optimized code will loop far fewer times than the code generated without the -O2 flag. Thus, the optimizer takes a lot of time (and some additional memory) preparing the code at compile-time in order to save a lot of time at run-time.